

CS 130 Project 2: User Programs Design Document

Tianyi Zhang

2018533074

zhangty2@shanghaitech.edu.cn

Haoran Dang

2018533259

danghr@shanghaitech.edu.cn

Derun Li

2018533152

lidr@shanghaitech.edu.cn

0 PRELIMINARIES

0.1 Preliminary Comments

No preliminary comment for this project.

0.2 References

- Pintos Guide by Stephen Tsung-Han Sher: https://static1.squarespace.com/static/5b18aa0955b02c1de94e4412/t/5e1bb4809e4b0a78012be132/1578874001361/Sher%282016%29_Pintos_Guide

1 ARGUMENT PASSING

1.1 Data Structures

1.1.1 *Copy here the declaration of each new or changed 'struct' or 'struct' member, global or static variable, 'typedef', or enumeration. Identify the purpose of each in 25 words or less. No modified data structure in this part.*

1.2 Algorithms

1.2.1 *Briefly describe how you implemented argument parsing. How do you arrange for the elements of argv[] to be in the right order? How do you avoid overflowing the stack page?* In function process_execute() of userprog/process.c, we create the new thread with the thread name picked by our newly-added function find_exec_name(), which picks the name of the executable from the file_name so that the name of the thread is correct.

Then in function setup_stack() of userprog/process.c, we add two pointers to store argc and argv after calling our newly-added find_args() in load() to separate them, which ensures that they are arranged in correct order. Then we push these arguments in correct places according to Section 3.5.1 of Pintos Document.

Number of args[] is limited by MAX_ARGS and MAX_CMD_LENGTH, which ensures that the stack is not overflowed.

1.3 Rationale

1.3.1 *Why does Pintos implement strtok_r() but not strtok()?* strtok() does not save the remaining part of the process but strtok_r() saves it in *save_ptr, which keeps the safety of the operation when multiple thread is accessing the same string.

1.3.2 *In Pintos, the kernel separates commands into a executable name and arguments. In Unix-like systems, the shell does this separation. Identify at least two advantages of the Unix approach.*

- It is easier to identify the file name instead of fetching from the command so that time can be saved when finding and opening the executable.

- If the name of the executable file contains special characters (like spaces), the kernel can also handle it easily.

2 SYSTEM CALLS

2.1 Data Structures

2.1.1 *Copy here the declaration of each new or changed 'struct' or 'struct' member, global or static variable, 'typedef', or enumeration. Identify the purpose of each in 25 words or less.*

- In file threads/threads.h

```
/* List element of waiting child */
struct waiting_sema
{
    /* List element */
    struct list_elem elem;
    /* TID of children */
    int child_tid;
    /* Semaphore corresponding to thread
       TID */
    struct semaphore sema;
};

struct thread
{
    ...
    /* Exit status of the thread */
    int exit_status;
    /* Next file descriptor */
    int next_fd;
    /* List of opened files */
    struct list opened_files;
    /* Pointer to the executable of
       thecurrent thread */
    struct file *executing_file;
    ...
    /* List used to store its child threads */
    struct list child_threads_list;
    /* List element of parent's
       child_threads_list */
    struct list_elem child_elem;
    /* List of waiting_sema of waiting
       threads */
    struct list waiting;
    /* Status of exited */
    bool is_exited;
    /* Status of being waited */
    bool is_waited;
    /* Status of waiting others */
};
```

```

    bool is_waiting;
    /* Pointer to parent thread */
    struct thread* parent_thread;
}

/* File operation lock */
struct lock file_lock;

```

- In file `userprog/syscall.c`

```

/* Interrupt handler wrapper functions */
static int (*syscall_handler_wrapper[20])
    (struct intr_frame *);

```
- In file `userprog/process.c`

```

/* Save exit status of the threads */
static int exit_status[MAX_THREADS] = {-1};
static bool is_exited[MAX_THREADS] = {false};

```

2.1.2 Describe how file descriptors are associated with open files. Are file descriptors unique within the entire OS or just within a single process? A fixed integer `next_fd` is saved in each thread. When opening a new file, if it exists, `allocate_fd()` will be called and it will return `next_fd++` whose value will be assigned to the opened file.

The file descriptors are unique within just a single process.

2.2 Algorithms

2.2.1 Describe your code for reading and writing user data from the kernel. Before accessing the data, we use a function `is_valid_addr()` which calls `is_user_vaddr()` and `pagedir_get_page()` to validate the user memory before dereferencing the pointers or values. Note that for each validation 4 bytes are needed to validate (and we can validate the front and the back) as they are pointers which save memory addresses.

When reading/writing these data, we dereference these pointers so that the corresponding data can be modified by the kernel.

2.2.2 Suppose a system call causes a full page (4,096 bytes) of data to be copied from user space into the kernel. What is the least and the greatest possible number of inspections of the page table (e.g. calls to `pagedir_get_page()`) that might result? What about for a system call that only copies 2 bytes of data? Is there room for improvement in these numbers, and how much? If the memory is contiguous, and we find the head of the page at the first inspection, then the number is 1. But if the memory is not continuous, the largest number will be 4096. For two bytes, the minimum number is 1, and the maximum number is 2 like above.

2.2.3 Briefly describe your implementation of the "wait" system call and how it interacts with process termination. We implement a `curr_child_list` to store the child process and find the process according to the `tid`, then we check whether the child process is waiting. If it is waiting then we return -1, otherwise we will push back the `waiting_sema` and block the parent thread until it is woken up by `process_exit` of its child thread.

When a thread is exiting, it will save its `exit_status` in the static array and change the `is_exited` to true. The `process_wait()` will then return the value of `exit_status[TID]` to `syscall_wait()`, which will save the status in the return value register.

2.2.4 Any access to user program memory at a user-specified address can fail due to a bad pointer value. Such accesses must cause the process to be terminated. System calls are fraught with such accesses, e.g. a "write" system call requires reading the system call number from the user stack, then each of the call's three arguments, then an arbitrary amount of user memory, and any of these can fail at any point. This poses a design and error-handling problem: how do you best avoid obscuring the primary function of code in a morass of error-handling? Furthermore, when an error is detected, how do you ensure that all temporarily allocated resources (locks, buffers, etc.) are freed? In a few paragraphs, describe the strategy or strategies you adopted for managing these issues. Give an example. We will check the last byte for the corresponding parameter. When we check the string we cannot get the length of the string. So we just check the char one by one until we find the EOF. And we will check the pointer stored in the stack which will terminate our program. If there is an error, we will terminate the process and then free the page.

2.3 Synchronization

2.3.1 The "exec" system call returns -1 if loading the new executable fails, so it cannot return before the new executable has completed loading. How does your code ensure this? How is the load success/failure status passed back to the thread that calls "exec"? We use a lock to synchronize the loading file part. Every time when we load a file we will acquire the only `file_lock` and when finish the load we will release the `file_lock`. We use the success to store whether we load the file or not.

2.3.2 Consider parent process *P* with child process *C*. How do you ensure proper synchronization and avoid race conditions when *P* calls `wait(C)` before *C* exits? After *C* exits? How do you ensure that all resources are freed in each case? How about when *P* terminates without waiting, before *C* exits? After *C* exits? Are there any special cases? We used two static arrays `is_exited` and `exit_status` to save the exit status. When a thread exits, it will set its `is_exited[TID]` to true and save its exit status to `exit_status[TID]`, so we can learn whether the thread is exited. If it is not exited, it works as stated in section 2.2.3; if it is exited, we straightly return the value from `exit_status`.

2.4 Rationale

2.4.1 Why did you choose to implement access to user memory from the kernel in the way that you did? In this way, dereferencing parameters and executing system calls are separated. Thus, we can validate the memory access before we actually run into system call codes. By the way, such a code is neat and easier to debug.

2.4.2 What advantages or disadvantages can you see to your design for file descriptors?

- Advantages: Each file has a unique file descriptor in the whole OS so it avoids most possible collisions.
- Disadvantages: The opened file is limited during a power cycle, since the maximum value of `int` is limited and file descriptors are never reallocated.

2.4.3 The default `tid_t` to `pid_t` mapping is the identity mapping. If you changed it, what advantages are there to your approach? We did not change it.

3 SURVEY QUESTIONS

In your opinion, was this assignment, or any one of the two problems in it, too easy or too hard? Did it take too long or too little time? Test-case multi-oom is really confusing. `backtrace` function of GDB is not easy to use in multi-process projects (like this one).

Did you find that working on a particular part of the assignment gave you greater insight into some aspect of OS design? Implementing argument passing and dereferencing parameters of system calls offers a insight to how process get and pass their parameters.

Is there some particular fact or hint we should give students in future quarters to help them solve the problems? Conversely, did you find any of our guidance to be misleading? Not really.

Do you have any suggestions for the TAs to more effectively assist students, either for future quarters or the remaining projects? Not yet.

Any other comments? Not yet.

CONTRIBUTORS

Task		Tianyi Zhang	Haoran Dang	Derun Li
Task 1 Argument Passing	Concept	✓	✓	✓
	Implementation	✓	✓	✓
	Debugging		✓	
	Design Document		✓	
Task 2 System Calls	Concept	✓	✓	✓
	Implementation	✓	✓	✓
	Debugging	✓	✓	✓
	Design Document	✓	✓	✓